

```
In [1]: # This Python 3 environment comes with many helpful analytics libraries installed
# It is defined by the kaggle/python Docker image: https://github.com/kaggle/docker-python
# For example, here's several helpful packages to load

import numpy as np # linear algebra
import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)

# Input data files are available in the read-only "../input/" directory
# For example, running this (by clicking run or pressing Shift+Enter) will list the files in the input directory

import os
for dirname, _, filenames in os.walk('/kaggle/input'):
    for filename in filenames:
        print(os.path.join(dirname, filename))

# You can write up to 20GB to the current directory (/kaggle/working/) that gets uploaded to S3
# You can also write temporary files to /kaggle/temp/, but they won't be saved outside this notebook
```

```
In [1]: import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import pandas as pd
import random
```

```
In [2]: pip install --upgrade protobuf==4.25.*
```

```
Requirement already satisfied: protobuf==4.25.* in c:\users\mohes\appdata\local\programs\python\python313\lib\site-packages (4.25.8)
Note: you may need to restart the kernel to use updated packages.
```

```
[notice] A new release of pip is available: 25.2 -> 25.3
[notice] To update, run: python.exe -m pip install --upgrade pip
```

```
In [4]: #%%pip install keras
import keras
```

```
In [5]: #%%pip install keras
import keras
from keras.models import Sequential
from keras.layers import Conv2D, Flatten, Dense, MaxPooling2D, Dropout
from sklearn.metrics import accuracy_score
```

```
In [57]: #%%pip install ipywidgets
# %%pip install tqdm
import ipywidgets as widgets
import io
from PIL import Image
import tqdm
from sklearn.model_selection import train_test_split
import cv2
from sklearn.utils import shuffle
import tensorflow as tf
```

```
In [10]: import os
import cv2
import numpy as np

X_train, Y_train = [], []
```

```

X_test, Y_test = [], []

image_size = 150
labels = ['glioma_tumor', 'meningioma_tumor', 'no_tumor', 'pituitary_tumor']

# Base paths for training and testing
train_base_path = r"C:\Users\mahes\OneDrive - United Nations\Data Science course"
test_base_path = r"C:\Users\mahes\OneDrive - United Nations\Data Science course"

# Load training data
for label in labels:
    folderPath = os.path.join(train_base_path, label)
    for filename in os.listdir(folderPath):
        img_path = os.path.join(folderPath, filename)
        img = cv2.imread(img_path)
        if img is not None:
            img = cv2.resize(img, (image_size, image_size))
            X_train.append(img)
            Y_train.append(label)

# Load testing data
for label in labels:
    folderPath = os.path.join(test_base_path, label)
    for filename in os.listdir(folderPath):
        img_path = os.path.join(folderPath, filename)
        img = cv2.imread(img_path)
        if img is not None:
            img = cv2.resize(img, (image_size, image_size))
            X_test.append(img)
            Y_test.append(label)

# Convert to numpy arrays
X_train = np.array(X_train)
Y_train = np.array(Y_train)
X_test = np.array(X_test)
Y_test = np.array(Y_test)

print(f"Training data: {X_train.shape}, Testing data: {X_test.shape}")

```

Training data: (2870, 150, 150, 3), Testing data: (394, 150, 150, 3)

In [11]: #  Basic dataset info

```

print("Total images:", len(X_train))
print("Image shape:", X_train[0].shape)
print("Unique labels:", np.unique(Y_train))

```

Total images: 2870  
Image shape: (150, 150, 3)  
Unique labels: ['glioma\_tumor' 'meningioma\_tumor' 'no\_tumor' 'pituitary\_tumor']

In [12]: #  Convert Labels to DataFrame for easy analysis

```

df = pd.DataFrame(Y_train, columns=['Tumor_Type'])
print("\nClass Distribution:")
print(df['Tumor_Type'].value_counts())

```

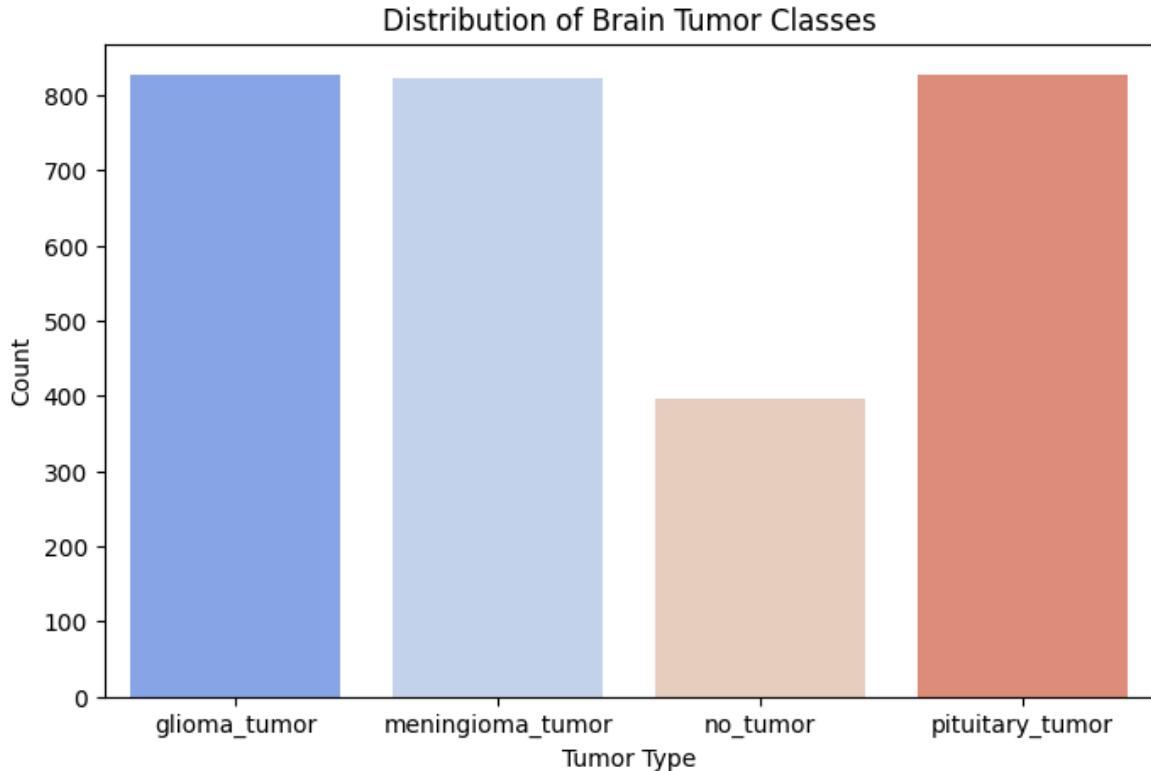
```
Class Distribution:  
Tumor_Type  
pituitary_tumor      827  
glioma_tumor        826  
meningioma_tumor    822  
no_tumor            395  
Name: count, dtype: int64
```

```
In [13]: # Plot class distribution  
plt.figure(figsize=(8,5))  
sns.countplot(x='Tumor_Type', data=df, palette='coolwarm')  
plt.title('Distribution of Brain Tumor Classes')  
plt.xlabel('Tumor Type')  
plt.ylabel('Count')  
plt.show()
```

C:\Users\mahes\AppData\Local\Temp\ipykernel\_9992\1265809927.py:3: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v 0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

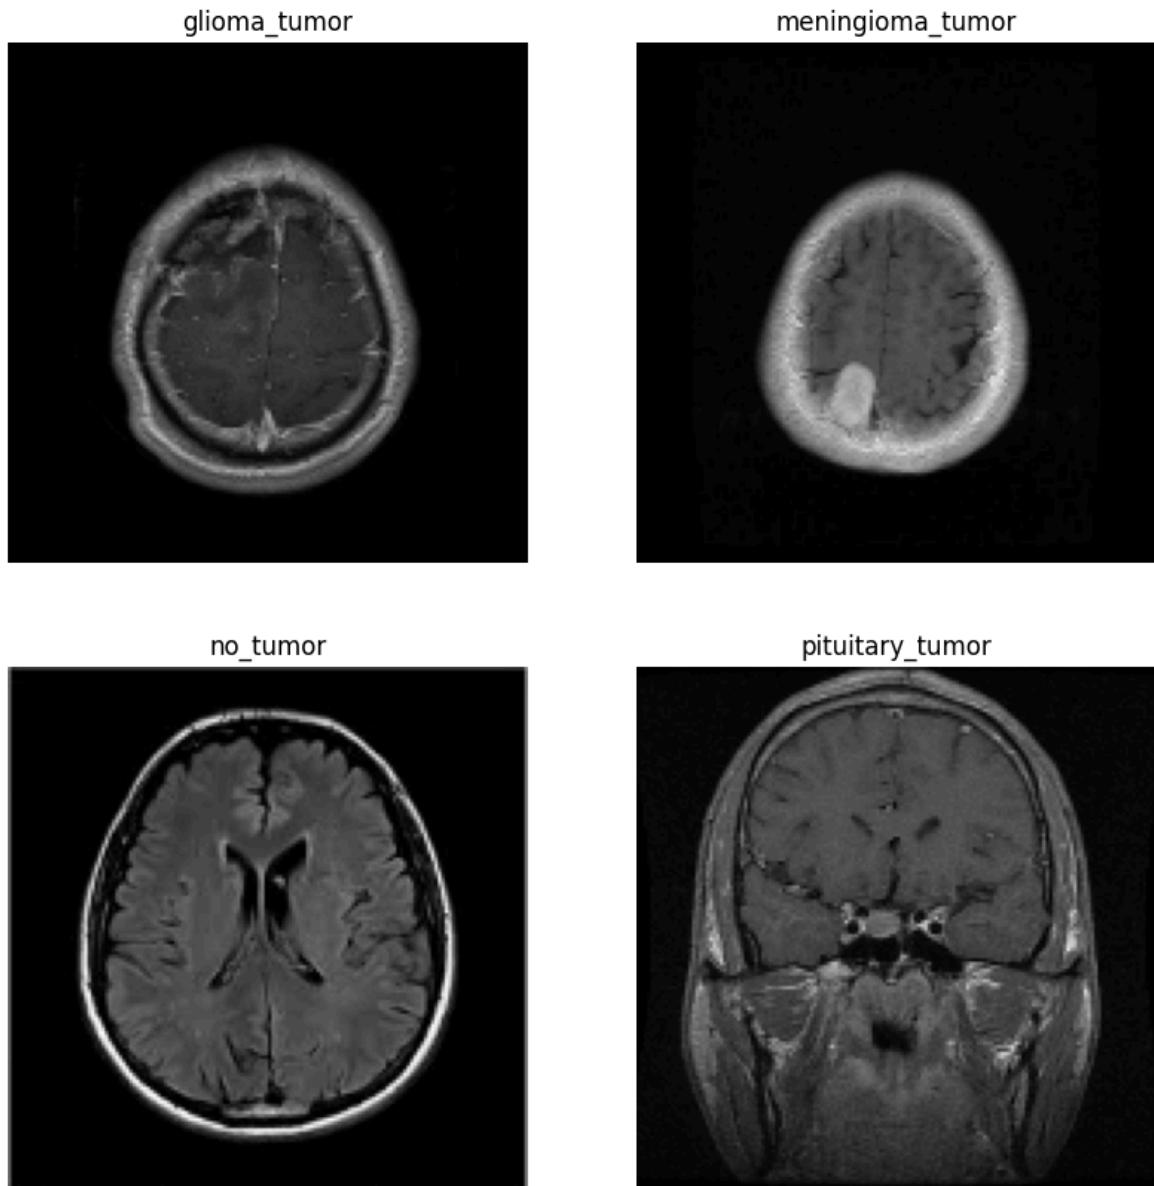
```
sns.countplot(x='Tumor_Type', data=df, palette='coolwarm')
```



```
In [14]: plt.figure(figsize=(10,10))  
for i, label in enumerate(labels):  
    plt.subplot(2,2,i+1)  
    # Find the indices that match the given label  
    indices = np.where(Y_train == label)[0]  
    if len(indices) > 0:  
        sample = random.choice(indices)  
        plt.imshow(cv2.cvtColor(X_train[sample], cv2.COLOR_BGR2RGB))  
        plt.title(label)  
        plt.axis('off')  
    else:  
        plt.text(0.3, 0.5, f"No images for {label}", fontsize=12)
```

```
plt.axis('off')
plt.suptitle("Random Samples from Each Brain Tumor Class", fontsize=16)
plt.show()
```

## Random Samples from Each Brain Tumor Class



```
In [16]: X_train,Y_train = shuffle(X_train,Y_train,random_state=0)
X_train.shape
```

```
Out[16]: (2870, 150, 150, 3)
```

```
In [17]: X_train,X_test,y_train,y_test = train_test_split(X_train,Y_train,test_size=0.1,r
```

```
In [18]: y_train_new = []
for i in y_train:
    y_train_new.append(labels.index(i))
y_train=y_train_new
y_train = tf.keras.utils.to_categorical(y_train)

y_test_new = []
for i in y_test:
    y_test_new.append(labels.index(i))
```

```
y_test=y_test_new  
y_test = tf.keras.utils.to_categorical(y_test)
```

```
In [19]: model = Sequential()  
model.add(Conv2D(32,(3,3),activation = 'relu',input_shape=(150,150,3)))  
model.add(Conv2D(64,(3,3),activation='relu'))  
model.add(MaxPooling2D(2,2))  
model.add(Dropout(0.3))  
model.add(Conv2D(64,(3,3),activation='relu'))  
model.add(Conv2D(64,(3,3),activation='relu'))  
model.add(Dropout(0.3))  
model.add(MaxPooling2D(2,2))  
model.add(Dropout(0.3))  
model.add(Conv2D(128,(3,3),activation='relu'))  
model.add(Conv2D(128,(3,3),activation='relu'))  
model.add(Conv2D(128,(3,3),activation='relu'))  
model.add(MaxPooling2D(2,2))  
model.add(Dropout(0.3))  
model.add(Conv2D(128,(3,3),activation='relu'))  
model.add(Conv2D(256,(3,3),activation='relu'))  
model.add(MaxPooling2D(2,2))  
model.add(Dropout(0.3))  
model.add(Flatten())  
model.add(Dense(512,activation = 'relu'))  
model.add(Dense(512,activation = 'relu'))  
model.add(Dropout(0.3))  
model.add(Dense(4,activation='softmax'))
```

```
c:\Users\mehes\AppData\Local\Programs\Python\Python313\Lib\site-packages\keras\s  
c\layers\convolutional\base_conv.py:113: UserWarning: Do not pass an `input_shape  
`/`input_dim` argument to a layer. When using Sequential models, prefer using an  
`Input(shape)` object as the first layer in the model instead.  
    super().__init__(activity_regularizer=activity_regularizer, **kwargs)
```

```
In [20]: model.summary()
```

```
Model: "sequential"
```

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 148, 148, 32)	896
conv2d_1 (Conv2D)	(None, 146, 146, 64)	18,496
max_pooling2d (MaxPooling2D)	(None, 73, 73, 64)	0
dropout (Dropout)	(None, 73, 73, 64)	0
conv2d_2 (Conv2D)	(None, 71, 71, 64)	36,928
conv2d_3 (Conv2D)	(None, 69, 69, 64)	36,928
dropout_1 (Dropout)	(None, 69, 69, 64)	0
max_pooling2d_1 (MaxPooling2D)	(None, 34, 34, 64)	0
dropout_2 (Dropout)	(None, 34, 34, 64)	0
conv2d_4 (Conv2D)	(None, 32, 32, 128)	73,856
conv2d_5 (Conv2D)	(None, 30, 30, 128)	147,584
conv2d_6 (Conv2D)	(None, 28, 28, 128)	147,584
max_pooling2d_2 (MaxPooling2D)	(None, 14, 14, 128)	0
dropout_3 (Dropout)	(None, 14, 14, 128)	0
conv2d_7 (Conv2D)	(None, 12, 12, 128)	147,584
conv2d_8 (Conv2D)	(None, 10, 10, 256)	295,168
max_pooling2d_3 (MaxPooling2D)	(None, 5, 5, 256)	0
dropout_4 (Dropout)	(None, 5, 5, 256)	0
flatten (Flatten)	(None, 6400)	0
dense (Dense)	(None, 512)	3,277,312
dense_1 (Dense)	(None, 512)	262,656
dropout_5 (Dropout)	(None, 512)	0
dense_2 (Dense)	(None, 4)	2,052

Total params: 4,447,044 (16.96 MB)

Trainable params: 4,447,044 (16.96 MB)

Non-trainable params: 0 (0.00 B)

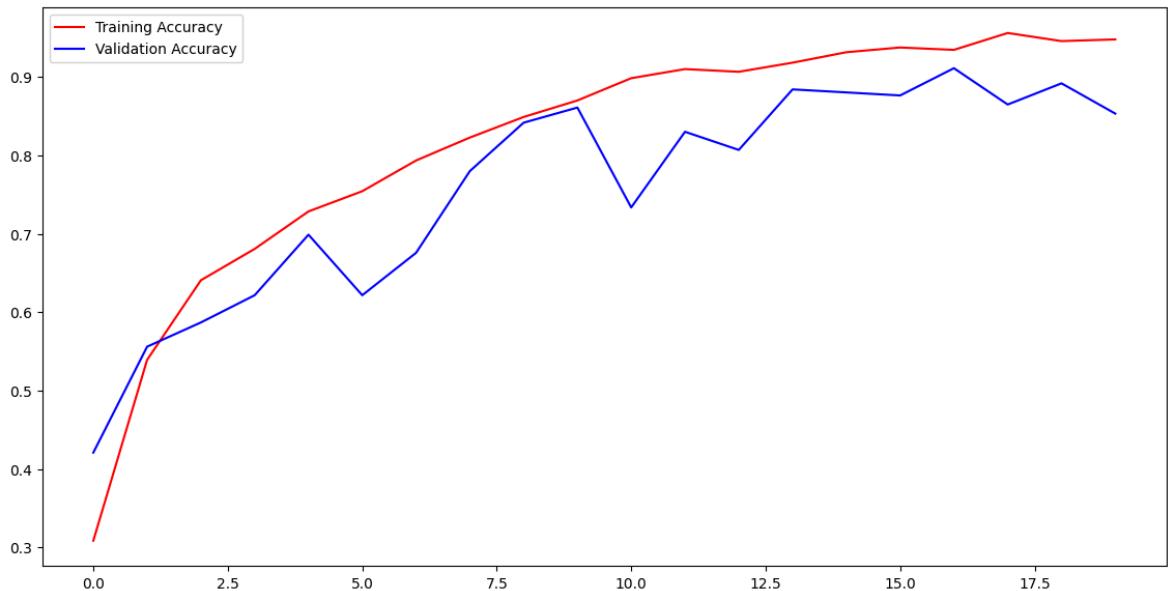
```
In [21]: model.compile(loss='categorical_crossentropy',optimizer='Adam',metrics=['accuracy'])

In [22]: history = model.fit(X_train,y_train,epochs=20,validation_split=0.1)
```

Epoch 1/20  
73/73 29s 375ms/step - accuracy: 0.3085 - loss: 1.6501 - val\_accuracy: 0.4208 - val\_loss: 1.3252  
Epoch 2/20  
73/73 31s 423ms/step - accuracy: 0.5392 - loss: 1.0609 - val\_accuracy: 0.5560 - val\_loss: 0.9304  
Epoch 3/20  
73/73 32s 443ms/step - accuracy: 0.6407 - loss: 0.8359 - val\_accuracy: 0.5869 - val\_loss: 0.9341  
Epoch 4/20  
73/73 33s 454ms/step - accuracy: 0.6807 - loss: 0.7407 - val\_accuracy: 0.6216 - val\_loss: 0.8933  
Epoch 5/20  
73/73 87s 1s/step - accuracy: 0.7285 - loss: 0.6464 - val\_accuracy: 0.6988 - val\_loss: 0.7203  
Epoch 6/20  
73/73 42s 580ms/step - accuracy: 0.7543 - loss: 0.5723 - val\_accuracy: 0.6216 - val\_loss: 0.9280  
Epoch 7/20  
73/73 43s 592ms/step - accuracy: 0.7935 - loss: 0.4938 - val\_accuracy: 0.6757 - val\_loss: 0.7463  
Epoch 8/20  
73/73 48s 660ms/step - accuracy: 0.8227 - loss: 0.4327 - val\_accuracy: 0.7799 - val\_loss: 0.5349  
Epoch 9/20  
73/73 44s 601ms/step - accuracy: 0.8490 - loss: 0.3763 - val\_accuracy: 0.8417 - val\_loss: 0.4284  
Epoch 10/20  
73/73 43s 590ms/step - accuracy: 0.8701 - loss: 0.3328 - val\_accuracy: 0.8610 - val\_loss: 0.4081  
Epoch 11/20  
73/73 42s 570ms/step - accuracy: 0.8985 - loss: 0.2751 - val\_accuracy: 0.7336 - val\_loss: 0.5721  
Epoch 12/20  
73/73 40s 553ms/step - accuracy: 0.9101 - loss: 0.2435 - val\_accuracy: 0.8301 - val\_loss: 0.4672  
Epoch 13/20  
73/73 40s 546ms/step - accuracy: 0.9066 - loss: 0.2598 - val\_accuracy: 0.8069 - val\_loss: 0.5505  
Epoch 14/20  
73/73 40s 545ms/step - accuracy: 0.9182 - loss: 0.2187 - val\_accuracy: 0.8842 - val\_loss: 0.3876  
Epoch 15/20  
73/73 40s 547ms/step - accuracy: 0.9316 - loss: 0.2027 - val\_accuracy: 0.8803 - val\_loss: 0.3411  
Epoch 16/20  
73/73 40s 554ms/step - accuracy: 0.9376 - loss: 0.1771 - val\_accuracy: 0.8764 - val\_loss: 0.3671  
Epoch 17/20  
73/73 40s 543ms/step - accuracy: 0.9346 - loss: 0.1827 - val\_accuracy: 0.9112 - val\_loss: 0.2939  
Epoch 18/20  
73/73 41s 563ms/step - accuracy: 0.9561 - loss: 0.1321 - val\_accuracy: 0.8649 - val\_loss: 0.4702  
Epoch 19/20  
73/73 40s 548ms/step - accuracy: 0.9458 - loss: 0.1343 - val\_accuracy: 0.8919 - val\_loss: 0.3440  
Epoch 20/20  
73/73 39s 539ms/step - accuracy: 0.9479 - loss: 0.1509 - val\_accuracy: 0.8533 - val\_loss: 0.3917

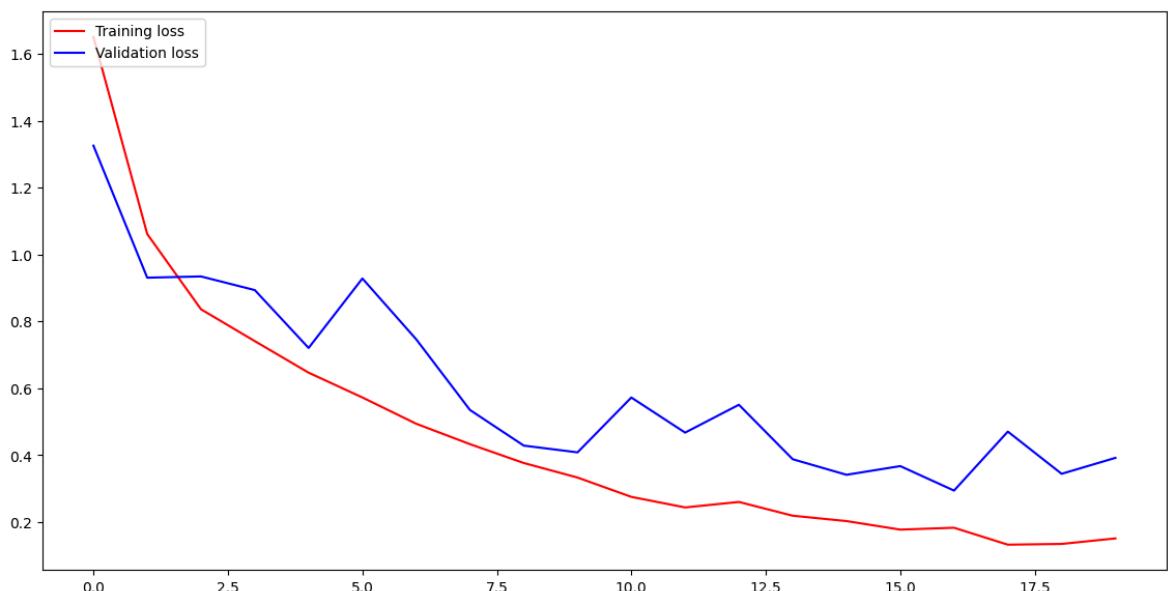
```
In [23]: #model.save('braintumor.h5')
```

```
acc = history.history['accuracy']
val_acc = history.history['val_accuracy']
epochs = range(len(acc))
fig = plt.figure(figsize=(14,7))
plt.plot(epochs, acc, 'r', label="Training Accuracy")
plt.plot(epochs, val_acc, 'b', label="Validation Accuracy")
plt.legend(loc='upper left')
plt.show()
```



```
In [24]: loss = history.history['loss']
val_loss = history.history['val_loss']
epochs = range(len(loss))
fig = plt.figure(figsize=(14,7))
```

```
plt.plot(epochs, loss, 'r', label="Training loss")
plt.plot(epochs, val_loss, 'b', label="Validation loss")
plt.legend(loc='upper left')
plt.show()
```



```
In [25]: test_loss, test_acc = model.evaluate(X_test, y_test)
print(f"Test Accuracy: {test_acc*100:.2f}%")
```

```
print(f"Test Loss: {test_loss:.4f}")

9/9 ----- 1s 92ms/step - accuracy: 0.8780 - loss: 0.3121
Test Accuracy: 87.80%
Test Loss: 0.3121
```

In [26]:

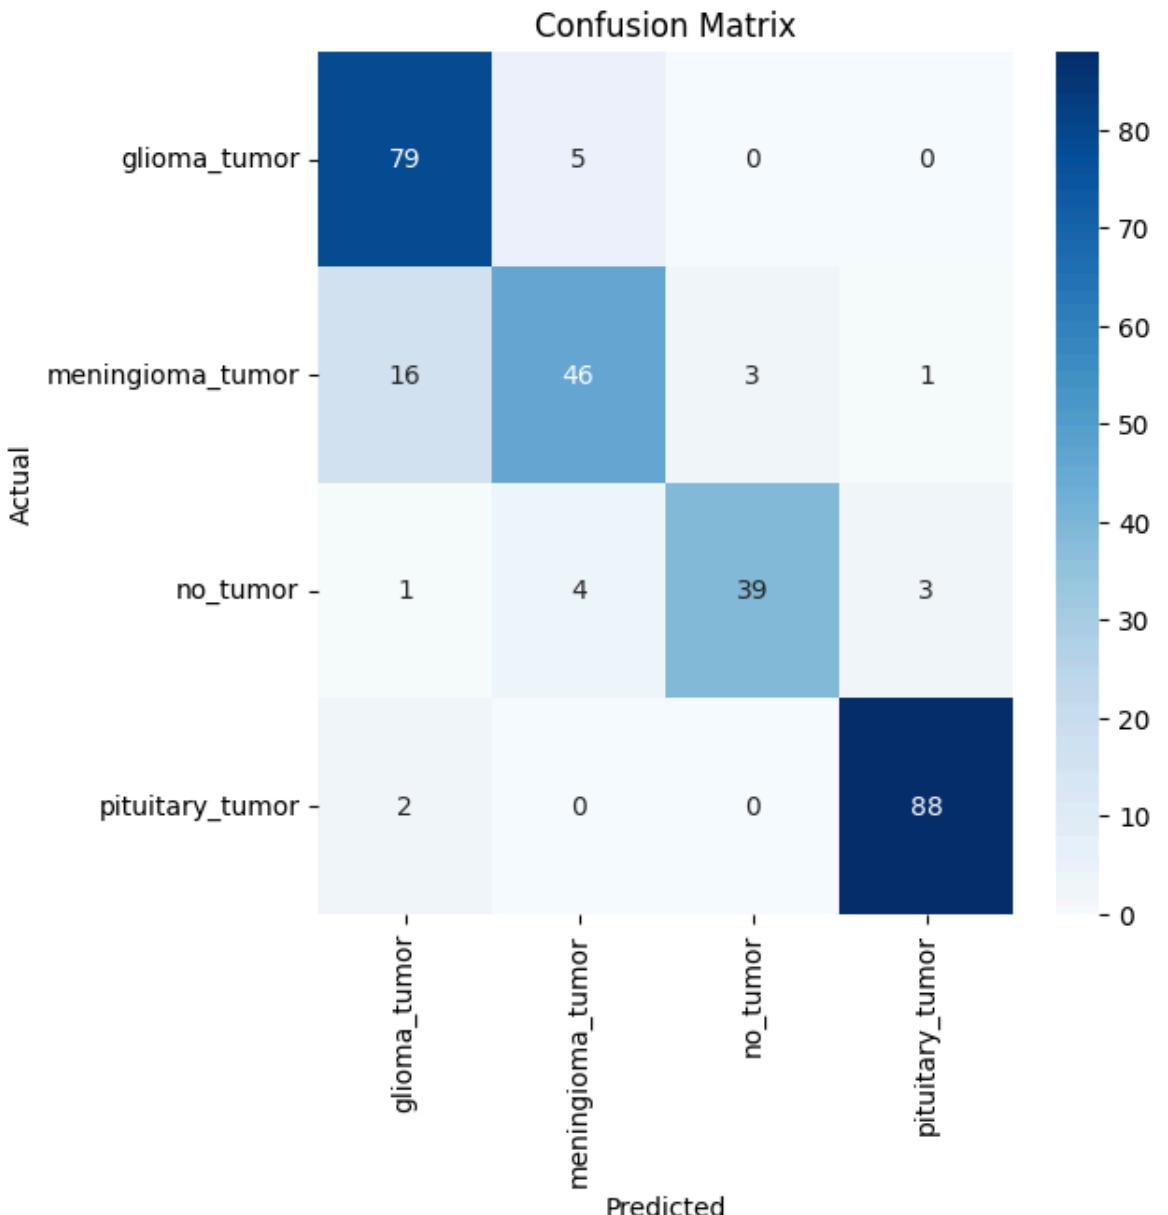
```
from sklearn.metrics import classification_report, confusion_matrix

# Predictions
y_pred = model.predict(X_test)
y_pred_classes = np.argmax(y_pred, axis=1)
y_true = np.argmax(y_test, axis=1)
```

```
9/9 ----- 1s 101ms/step
```

In [27]:

```
# Confusion matrix
cm = confusion_matrix(y_true, y_pred_classes)
plt.figure(figsize=(6,6))
sns.heatmap(cm, annot=True, fmt='d', xticklabels=labels, yticklabels=labels, cma
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.title("Confusion Matrix")
plt.show()
```



```
In [28]: # Classification report
print(classification_report(y_true, y_pred_classes, target_names=labels))
```

	precision	recall	f1-score	support
glioma_tumor	0.81	0.94	0.87	84
meningioma_tumor	0.84	0.70	0.76	66
no_tumor	0.93	0.83	0.88	47
pituitary_tumor	0.96	0.98	0.97	90
accuracy			0.88	287
macro avg	0.88	0.86	0.87	287
weighted avg	0.88	0.88	0.88	287

```
In [45]: import cv2
import numpy as np

img = cv2.imread(r"C:\Users\mahes\OneDrive - United Nations\Data Science course\pituitary_tumor_0.jpg")

if img is None:
    print("X Image not found. Check your file path.")
else:
    print("✓ Image loaded successfully.")

img = cv2.resize(img, (150, 150))
img_array = np.array(img)
img_array.shape
img_array = img_array.reshape(1, 150, 150, 3)
prediction = model.predict(img_array)
predicted_class = labels[np.argmax(prediction)]
print(f"The predicted class for the input image is: {predicted_class}")
```

✓ Image loaded successfully.  
1/1 ————— 0s 39ms/step  
The predicted class for the input image is: pituitary\_tumor  
1/1 ————— 0s 39ms/step  
The predicted class for the input image is: pituitary\_tumor

```
In [46]: img_array = img_array.reshape(1,150,150,3)
img_array.shape
```

Out[46]: (1, 150, 150, 3)

```
In [54]: import numpy as np
from tensorflow.keras.preprocessing import image

# Path to your image
img_path = (r"C:\Users\mahes\OneDrive - United Nations\Data Science course\Nares\pituitary_tumor_0.jpg")

if img is None:
    print("X Image not found. Check your file path.")
else:
    print("✓ Image loaded successfully.")
```

```

# Load image and resize to 150x150
img = image.load_img(img_path, target_size=(150, 150))

# Convert to numpy array
img_array = image.img_to_array(img)

# Add batch dimension (needed for prediction)
img_array = np.expand_dims(img_array, axis=0)

# Predict using your trained model
prediction = model.predict(img_array)
predicted_class = labels[np.argmax(prediction)]

print(f"The predicted class for the input image is: {predicted_class}")

```

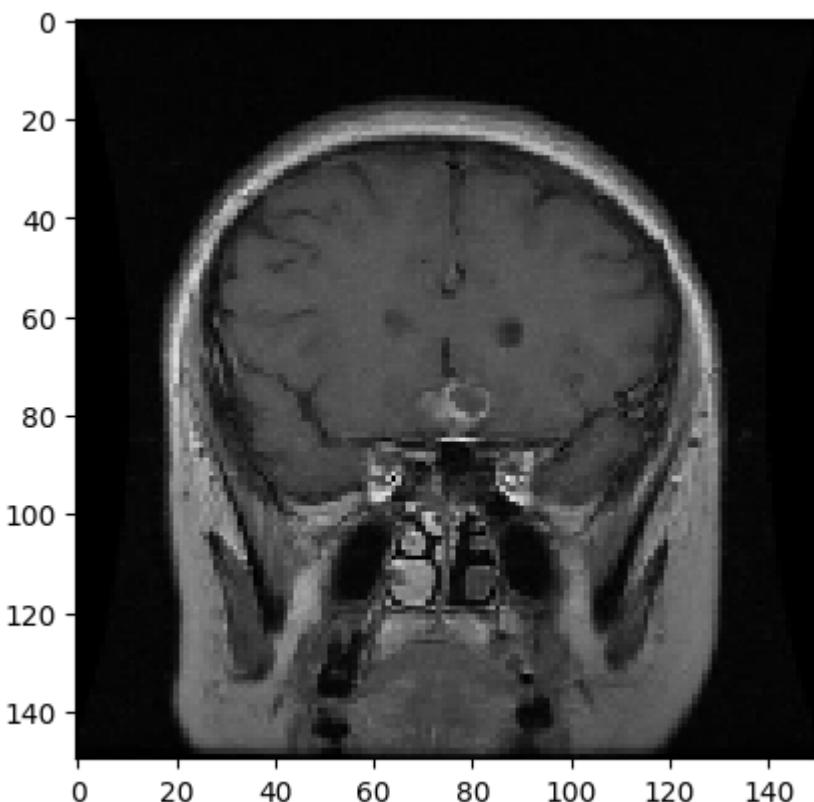
Image loaded successfully.  
 1/1 ————— 0s 86ms/step  
 1/1 ————— 0s 86ms/step  
 The predicted class for the input image is: pituitary\_tumor

In [55]:

```

from tensorflow.keras.preprocessing import image
img_path = (r"C:\Users\mahes\OneDrive - United Nations\Data Science course Nares"
plt.imshow(img, interpolation='nearest')
plt.show()

```



In [56]:

```

a=model.predict(img_array)
indices = a.argmax()
indices

```

1/1 ————— 0s 29ms/step

Out[56]:

```
np.int64(3)
```

In [ ]:

```

# In order to improve the model accuracy further, we can consider techniques such
# architectures or pre-trained models.

```

Hyperparameter	What it does
**Learning rate**	Step size for gradient descent
**Batch size**	Number of images processed before updating weights
**Number of epochs**	How many passes over the dataset
**Optimizer**	How weights are updated
**Dropout rate**	Prevents overfitting
**Number of layers / filters**	Model complexity

Out[58]:	'\n  Hyperparameter   What it does
	'\n  How to tune   \n  -----
	-----   -----   -----
	-----   \n  **Learning rate**
	'\n  Step size for gradient descent   Try smaller/larger rates
	(e.g., 0.001, 0.0005, 0.0001)   \n  **Batch size**   Number o
	f images processed before updating weights   Experiment (16, 32, 64)
	\n  **Number of epochs**   How many passes over the dataset
	'\n  Monitor validation loss to avoid overfitting   \n  **Optimizer**
	'\n  How weights are updated   Try `adam`, `sgd`, `rmsp
	rop`   \n  **Dropout rate**   Prevent
	s overfitting   Add 0.2 - 0.5 in dense layers
	\n  **Number of layers / filters**   Model complexity
	'\n  More layers/filters can improve accuracy, but can overfit   \n'

```
In [ ]: #Data Augmentation
"""
With MRI images, you often have limited data. Using augmentation artificially increases the size of your training set by applying various transformations to your images.

Flip images horizontally/vertically

Rotate slightly (10 - 20°)

Zoom, shift, shear

Adjust brightness/contrast

Example in Keras:

"""
from tensorflow.keras.preprocessing.image import ImageDataGenerator

train_datagen = ImageDataGenerator(
    rescale=1./255,
    rotation_range=15,
    width_shift_range=0.1,
    height_shift_range=0.1,
    shear_range=0.1,
    zoom_range=0.1,
    horizontal_flip=True,
    fill_mode='nearest'
)
# This helps the model generalize better, reducing overfitting and improving validation accuracy.
```

```
In [60]: # Regularization Techniques
```

```
from tensorflow.keras import regularizers
Dense(128, activation='relu', kernel_regularizer=regularizers.l2(0.001))
```

Out[60]: <Dense name=dense\_3, built=False>

In [61]: *# Early Stopping and Learning Rate Scheduler  
# Stop training when validation Loss stops improving to prevent overfitting.*

```
from tensorflow.keras.callbacks import EarlyStopping, ReduceLROnPlateau

early_stop = EarlyStopping(monitor='val_loss', patience=5, restore_best_weights=True)
reduce_lr = ReduceLROnPlateau(monitor='val_loss', factor=0.5, patience=3)
model.fit(X_train, y_train, epochs=50, validation_split=0.1, callbacks=[early_st
```

```
Epoch 1/50
73/73 _____ 28s 384ms/step - accuracy: 0.9574 - loss: 0.1213 - val
_accuracy: 0.8842 - val_loss: 0.4026 - learning_rate: 0.0010
Epoch 2/50
73/73 _____ 32s 433ms/step - accuracy: 0.9703 - loss: 0.0938 - val
_accuracy: 0.8185 - val_loss: 0.5745 - learning_rate: 0.0010
Epoch 3/50
73/73 _____ 32s 439ms/step - accuracy: 0.9634 - loss: 0.1029 - val
_accuracy: 0.9035 - val_loss: 0.3826 - learning_rate: 0.0010
Epoch 4/50
73/73 _____ 32s 443ms/step - accuracy: 0.9763 - loss: 0.0708 - val
_accuracy: 0.9305 - val_loss: 0.4757 - learning_rate: 0.0010
Epoch 5/50
73/73 _____ 32s 445ms/step - accuracy: 0.9686 - loss: 0.0873 - val
_accuracy: 0.8842 - val_loss: 0.3684 - learning_rate: 0.0010
Epoch 6/50
73/73 _____ 33s 454ms/step - accuracy: 0.9686 - loss: 0.0932 - val
_accuracy: 0.8533 - val_loss: 0.4579 - learning_rate: 0.0010
Epoch 7/50
73/73 _____ 33s 457ms/step - accuracy: 0.9660 - loss: 0.1035 - val
_accuracy: 0.9189 - val_loss: 0.4564 - learning_rate: 0.0010
Epoch 8/50
73/73 _____ 34s 469ms/step - accuracy: 0.9798 - loss: 0.0596 - val
_accuracy: 0.8880 - val_loss: 0.4832 - learning_rate: 0.0010
Epoch 9/50
73/73 _____ 35s 478ms/step - accuracy: 0.9793 - loss: 0.0511 - val
_accuracy: 0.8996 - val_loss: 0.3493 - learning_rate: 5.0000e-04
Epoch 10/50
73/73 _____ 36s 490ms/step - accuracy: 0.9910 - loss: 0.0292 - val
_accuracy: 0.9228 - val_loss: 0.3616 - learning_rate: 5.0000e-04
Epoch 11/50
73/73 _____ 37s 502ms/step - accuracy: 0.9884 - loss: 0.0426 - val
_accuracy: 0.9305 - val_loss: 0.3402 - learning_rate: 5.0000e-04
Epoch 12/50
73/73 _____ 37s 513ms/step - accuracy: 0.9948 - loss: 0.0155 - val
_accuracy: 0.9073 - val_loss: 0.4336 - learning_rate: 5.0000e-04
Epoch 13/50
73/73 _____ 39s 532ms/step - accuracy: 0.9927 - loss: 0.0191 - val
_accuracy: 0.9151 - val_loss: 0.3392 - learning_rate: 5.0000e-04
Epoch 14/50
73/73 _____ 39s 539ms/step - accuracy: 0.9957 - loss: 0.0142 - val
_accuracy: 0.9228 - val_loss: 0.4253 - learning_rate: 5.0000e-04
Epoch 15/50
73/73 _____ 40s 542ms/step - accuracy: 0.9935 - loss: 0.0229 - val
_accuracy: 0.9151 - val_loss: 0.3976 - learning_rate: 5.0000e-04
Epoch 16/50
73/73 _____ 40s 546ms/step - accuracy: 0.9957 - loss: 0.0110 - val
_accuracy: 0.9073 - val_loss: 0.3912 - learning_rate: 5.0000e-04
Epoch 17/50
73/73 _____ 41s 558ms/step - accuracy: 0.9978 - loss: 0.0057 - val
_accuracy: 0.9151 - val_loss: 0.4235 - learning_rate: 2.5000e-04
Epoch 18/50
73/73 _____ 41s 562ms/step - accuracy: 0.9978 - loss: 0.0075 - val
_accuracy: 0.9189 - val_loss: 0.4572 - learning_rate: 2.5000e-04
```

Out[61]: <keras.src.callbacks.history.History at 0x1f89ea29810>

Understanding each metric:

Epoch 17/50 & 18/50: In the above code, 17th and 18th training cycles out of a total of 50. Each epoch means my model has seen all training images once.

accuracy: 0.9978 ( $\approx$  99.78%) This is the model's accuracy on the training dataset. The model predicts the correct tumor class almost perfectly during training — excellent, but it might also hint at overfitting (too good on training, slightly worse on validation).

loss: 0.0057 → 0.0075 The training loss is extremely low — this means the model's predictions on training data are very close to the actual labels.

val\_accuracy: 0.9151 → 0.9189 ( $\approx$  91.9%) The accuracy on validation data, which represents new/unseen images. This is quite solid for medical image classification! A model with >90% validation accuracy is promising — especially if the data is real MRI scans.

val\_loss: 0.4235 → 0.4572 The validation loss is slowly increasing — this might be an early sign of overfitting, meaning the model memorizes training data patterns but struggles with unseen data.

learning\_rate: 2.5e-04 (0.00025) A relatively low learning rate — ideal for fine-tuning. It helps the model converge smoothly instead of overshooting the optimal weights.

In the next block, I am trying to improve validation accuracy further and avoid overfitting:

```
In [ ]: ## improve validation accuracy further and avoid overfitting:  
  
from tensorflow.keras.callbacks import EarlyStopping, ReduceLROnPlateau  
  
# Stop training early if validation loss doesn't improve for 5 epochs  
early_stop = EarlyStopping(  
    monitor='val_loss',  
    patience=5,  
    restore_best_weights=True  
)  
  
# Reduce Learning rate by 50% if validation loss doesn't improve for 3 epochs  
reduce_lr = ReduceLROnPlateau(  
    monitor='val_loss',  
    factor=0.5,  
    patience=3  
)  
  
# Train model with both callbacks  
model.fit(  
    X_train,  
    y_train,  
    epochs=50,  
    validation_split=0.1,  
    callbacks=[early_stop, reduce_lr])
```

```
Epoch 1/50
73/73 29s 397ms/step - accuracy: 0.9983 - loss: 0.0070 - val_accuracy: 0.9035 - val_loss: 0.3844 - learning_rate: 2.5000e-04
Epoch 2/50
73/73 34s 458ms/step - accuracy: 0.9978 - loss: 0.0077 - val_accuracy: 0.9151 - val_loss: 0.4287 - learning_rate: 2.5000e-04
Epoch 3/50
73/73 34s 468ms/step - accuracy: 0.9996 - loss: 0.0048 - val_accuracy: 0.9228 - val_loss: 0.4069 - learning_rate: 2.5000e-04
Epoch 4/50
73/73 33s 457ms/step - accuracy: 0.9983 - loss: 0.0036 - val_accuracy: 0.9266 - val_loss: 0.4282 - learning_rate: 2.5000e-04
Epoch 5/50
73/73 34s 460ms/step - accuracy: 0.9991 - loss: 0.0027 - val_accuracy: 0.9305 - val_loss: 0.4330 - learning_rate: 1.2500e-04
Epoch 6/50
73/73 35s 478ms/step - accuracy: 0.9983 - loss: 0.0056 - val_accuracy: 0.9305 - val_loss: 0.4442 - learning_rate: 1.2500e-04
Out[ ]: <keras.callbacks.history.History at 0x1f89ea29950>
```

What the above means:

Training accuracy: 0.9983 → this model is fitting almost perfectly to training data.

Validation accuracy: 0.9305 → That's excellent performance on unseen data (93%).

Validation loss: 0.4442 → Acceptable; slightly higher than training loss, meaning there's minor overfitting, but nothing alarming.

Learning rate: 1.25e-04 → The ReduceLROnPlateau callback has likely reduced the learning rate, helping fine-tune the model gradually.

Epoch: 6/50 → And EarlyStopping may stop earlier than 50 epochs once validation loss stops improving.

```
In [63]: # Let it train until early stopping triggers

test_loss, test_acc = model.evaluate(X_test, y_test)
print(f"Test Accuracy: {test_acc:.4f}, Test Loss: {test_loss:.4f}")

9/9 1s 93ms/step - accuracy: 0.9338 - loss: 0.2995
Test Accuracy: 0.9338, Test Loss: 0.2995
```

Check confusion matrix & classification report to confirm per-class performance:

```
In [64]: from sklearn.metrics import classification_report, confusion_matrix
y_pred = model.predict(X_test)
y_pred_classes = np.argmax(y_pred, axis=1)
y_true = np.argmax(y_test, axis=1)

print(classification_report(y_true, y_pred_classes))
print(confusion_matrix(y_true, y_pred_classes))
```

```
9/9 ━━━━━━ 1s 95ms/step
      precision    recall   f1-score   support
          0       0.88     0.94     0.91      84
          1       0.91     0.88     0.89      66
          2       0.96     0.94     0.95      47
          3       1.00     0.97     0.98      90
accuracy                      0.93      287
macro avg        0.94     0.93     0.93      287
weighted avg     0.94     0.93     0.93      287

[[79  5  0  0]
 [ 6 58  2  0]
 [ 2  1 44  0]
 [ 3  0  0 87]]
```

```
In [66]: model.save("brain_tumor_classifier_v1.keras")
```

```
In [67]: from tensorflow import keras
```

```
# Load the trained model
model = keras.models.load_model("brain_tumor_classifier_v1.keras")
print("✓ Model loaded successfully!")
```

✓ Model loaded successfully!

```
c:\Users\mahes\AppData\Local\Programs\Python\Python313\Lib\site-packages\keras\saving\saving_lib.py:797: UserWarning: Skipping variable loading for optimizer 'rmsprop', because it has 26 variables whereas the saved optimizer has 50 variables.
    saveable.load_own_variables(weights_store.get(inner_path))
```

```
In [68]: # Predict on a new MRI image
```

```
labels = ['glioma_tumor', 'meningioma_tumor', 'no_tumor', 'pituitary_tumor']
```

```
In [ ]: import numpy as np
from tensorflow.keras.preprocessing import image

# Path to the new MRI image
img_path = r"C:\path\to\your\test_image.jpg" ## Change this to your image path (


# Load and preprocess the image
img = image.load_img(img_path, target_size=(150, 150))
img_array = image.img_to_array(img)
img_array = np.expand_dims(img_array, axis=0) # Add batch dimension
img_array = img_array / 255.0 # Normalize if your training data was normalized

# Predict
prediction = model.predict(img_array)
predicted_class = labels[np.argmax(prediction)]

print(f"💡 The predicted tumor type is: {predicted_class}")
```

```
1/1 ━━━━━━ 0s 89ms/step
```

```
1/1 ━━━━━━ 0s 89ms/step
```

💡 The predicted tumor type is: glioma\_tumor

```
In [ ]: #%%pip install streamlit
import streamlit as st
```

```

import streamlit as st
import numpy as np
from tensorflow import keras
from tensorflow.keras.preprocessing import image
from PIL import Image

# Load the trained model
st.title("🧠 Brain Tumor Classification App")
st.write("Upload an MRI brain scan image to predict the tumor type.")

# Load model
@st.cache_resource
def load_model():
    model = keras.models.load_model("brain_tumor_classifier_v1.keras")
    return model

model = load_model()

# Labels
labels = ['glioma_tumor', 'meningioma_tumor', 'no_tumor', 'pituitary_tumor']

# File uploader
uploaded_file = st.file_uploader("Upload an MRI image", type=["jpg", "jpeg", "png"])

if uploaded_file is not None:
    # Display uploaded image
    image_display = Image.open(uploaded_file)
    st.image(image_display, caption="Uploaded MRI Image", use_column_width=True)

    # Preprocess the image
    img = image.load_img(uploaded_file, target_size=(150, 150))
    img_array = image.img_to_array(img)
    img_array = np.expand_dims(img_array, axis=0) / 255.0 # normalize

    # Predict
    prediction = model.predict(img_array)
    predicted_class = labels[np.argmax(prediction)]
    confidence = np.max(prediction)

    # Display results
    st.markdown(f"### 🎨 Predicted Tumor Type: **{predicted_class}**")
    st.markdown(f"### 🔳 Confidence: **{confidence:.2f}**")

```

Requirement already satisfied: streamlit in c:\users\mahes\appdata\local\programs\python\python313\lib\site-packages (1.49.1)  
Requirement already satisfied: altair!=5.4.0,!\_=5.4.1,<6,>=4.0 in c:\users\mahes\appdata\local\programs\python\python313\lib\site-packages (from streamlit) (5.5.0)  
Requirement already satisfied: blinker<2,>=1.5.0 in c:\users\mahes\appdata\local\programs\python\python313\lib\site-packages (from streamlit) (1.9.0)  
Requirement already satisfied: cachetools<7,>=4.0 in c:\users\mahes\appdata\local\programs\python\python313\lib\site-packages (from streamlit) (6.2.0)  
Requirement already satisfied: click<9,>=7.0 in c:\users\mahes\appdata\local\programs\python\python313\lib\site-packages (from streamlit) (8.1.8)  
Requirement already satisfied: numpy<3,>=1.23 in c:\users\mahes\appdata\local\programs\python\python313\lib\site-packages (from streamlit) (2.2.6)  
Requirement already satisfied: packaging<26,>=20 in c:\users\mahes\appdata\local\programs\python\python313\lib\site-packages (from streamlit) (25.0)  
Requirement already satisfied: pandas<3,>=1.4.0 in c:\users\mahes\appdata\local\programs\python\python313\lib\site-packages (from streamlit) (2.3.3)  
Requirement already satisfied: pillow<12,>=7.1.0 in c:\users\mahes\appdata\local\programs\python\python313\lib\site-packages (from streamlit) (11.3.0)  
Requirement already satisfied: protobuf<7,>=3.20 in c:\users\mahes\appdata\local\programs\python\python313\lib\site-packages (from streamlit) (6.33.0)  
Requirement already satisfied: pyarrow>=7.0 in c:\users\mahes\appdata\local\programs\python\python313\lib\site-packages (from streamlit) (21.0.0)  
Requirement already satisfied: requests<3,>=2.27 in c:\users\mahes\appdata\local\programs\python\python313\lib\site-packages (from streamlit) (2.32.5)  
Requirement already satisfied: tenacity<10,>=8.1.0 in c:\users\mahes\appdata\local\programs\python\python313\lib\site-packages (from streamlit) (9.1.2)  
Requirement already satisfied: toml<2,>=0.10.1 in c:\users\mahes\appdata\local\programs\python\python313\lib\site-packages (from streamlit) (0.10.2)  
Requirement already satisfied: typing-extensions<5,>=4.4.0 in c:\users\mahes\appdata\local\programs\python\python313\lib\site-packages (from streamlit) (4.15.0)  
Requirement already satisfied: watchdog<7,>=2.1.5 in c:\users\mahes\appdata\local\programs\python\python313\lib\site-packages (from streamlit) (6.0.0)  
Requirement already satisfied: gitpython!=3.1.19,<4,>=3.0.7 in c:\users\mahes\appdata\local\programs\python\python313\lib\site-packages (from streamlit) (3.1.45)  
Requirement already satisfied: pydeck<1,>=0.8.0b4 in c:\users\mahes\appdata\local\programs\python\python313\lib\site-packages (from streamlit) (0.9.1)  
Requirement already satisfied: tornado!=6.5.0,<7,>=6.0.3 in c:\users\mahes\appdata\local\programs\python\python313\lib\site-packages (from streamlit) (6.5.2)  
Requirement already satisfied: jinja2 in c:\users\mahes\appdata\local\programs\python\python313\lib\site-packages (from altair!=5.4.0,!\_=5.4.1,<6,>=4.0->streamlit) (3.1.6)  
Requirement already satisfied: jsonschema>=3.0 in c:\users\mahes\appdata\local\programs\python\python313\lib\site-packages (from altair!=5.4.0,!\_=5.4.1,<6,>=4.0->streamlit) (4.25.1)  
Requirement already satisfied: narwhal>=1.14.2 in c:\users\mahes\appdata\local\programs\python\python313\lib\site-packages (from altair!=5.4.0,!\_=5.4.1,<6,>=4.0->streamlit) (2.3.0)  
Requirement already satisfied: colorama in c:\users\mahes\appdata\local\programs\python\python313\lib\site-packages (from click<9,>=7.0->streamlit) (0.4.6)  
Requirement already satisfied: gitdb<5,>=4.0.1 in c:\users\mahes\appdata\local\programs\python\python313\lib\site-packages (from gitpython!=3.1.19,<4,>=3.0.7->streamlit) (4.0.12)  
Requirement already satisfied: smmap<6,>=3.0.1 in c:\users\mahes\appdata\local\programs\python\python313\lib\site-packages (from gitdb<5,>=4.0.1->gitpython!=3.1.19,<4,>=3.0.7->streamlit) (5.0.2)  
Requirement already satisfied: python-dateutil>=2.8.2 in c:\users\mahes\appdata\local\programs\python\python313\lib\site-packages (from pandas<3,>=1.4.0->streamlit) (2.9.0.post0)  
Requirement already satisfied: pytz>=2020.1 in c:\users\mahes\appdata\local\programs\python\python313\lib\site-packages (from pandas<3,>=1.4.0->streamlit) (2025.0)

2)  
Requirement already satisfied: tzdata>=2022.7 in c:\users\mahes\appdata\local\programs\python\python313\lib\site-packages (from pandas<3,>=1.4.0->streamlit) (2025.2)  
Requirement already satisfied: charset\_normalizer<4,>=2 in c:\users\mahes\appdata\local\programs\python\python313\lib\site-packages (from requests<3,>=2.27->streamlit) (3.4.3)  
Requirement already satisfied: idna<4,>=2.5 in c:\users\mahes\appdata\local\programs\python\python313\lib\site-packages (from requests<3,>=2.27->streamlit) (3.10)  
Requirement already satisfied: urllib3<3,>=1.21.1 in c:\users\mahes\appdata\local\programs\python\python313\lib\site-packages (from requests<3,>=2.27->streamlit) (2.5.0)  
Requirement already satisfied: certifi>=2017.4.17 in c:\users\mahes\appdata\local\programs\python\python313\lib\site-packages (from requests<3,>=2.27->streamlit) (2025.8.3)  
Requirement already satisfied: MarkupSafe>=2.0 in c:\users\mahes\appdata\local\programs\python\python313\lib\site-packages (from jinja2->altair!=5.4.0,!>5.4.1,<6,>=4.0->streamlit) (3.0.2)  
Requirement already satisfied: attrs>=22.2.0 in c:\users\mahes\appdata\local\programs\python\python313\lib\site-packages (from jsonschema>=3.0->altair!=5.4.0,!>5.4.1,<6,>=4.0->streamlit) (25.3.0)  
Requirement already satisfied: jsonschema-specifications>=2023.03.6 in c:\users\mahes\appdata\local\programs\python\python313\lib\site-packages (from jsonschema>=3.0->altair!=5.4.0,!>5.4.1,<6,>=4.0->streamlit) (2025.4.1)  
Requirement already satisfied: referencing>=0.28.4 in c:\users\mahes\appdata\local\programs\python\python313\lib\site-packages (from jsonschema>=3.0->altair!=5.4.0,!>5.4.1,<6,>=4.0->streamlit) (0.36.2)  
Requirement already satisfied: rpds-py>=0.7.1 in c:\users\mahes\appdata\local\programs\python\python313\lib\site-packages (from jsonschema>=3.0->altair!=5.4.0,!>5.4.1,<6,>=4.0->streamlit) (0.27.1)  
Requirement already satisfied: six>=1.5 in c:\users\mahes\appdata\local\programs\python\python313\lib\site-packages (from python-dateutil>=2.8.2->pandas<3,>=1.4.0->streamlit) (1.17.0)

Note: you may need to restart the kernel to use updated packages.

```
[notice] A new release of pip is available: 25.2 -> 25.3
[notice] To update, run: python.exe -m pip install --upgrade pip
2025-11-09 17:12:38.057 WARNING streamlit.runtime.scriptrunner_utils.script_run_context: Thread 'MainThread': missing ScriptRunContext! This warning can be ignored when running in bare mode.
2025-11-09 17:12:38.401
Warning: to view this Streamlit app on a browser, run it with the following command:

    streamlit run C:\Users\mahes\AppData\Roaming\Python\Python313\site-packages\ipykernel_launcher.py [ARGUMENTS]
2025-11-09 17:12:38.402 Thread 'MainThread': missing ScriptRunContext! This warning can be ignored when running in bare mode.
2025-11-09 17:12:38.403 Thread 'MainThread': missing ScriptRunContext! This warning can be ignored when running in bare mode.
2025-11-09 17:12:38.403 Thread 'MainThread': missing ScriptRunContext! This warning can be ignored when running in bare mode.
2025-11-09 17:12:38.404 Thread 'MainThread': missing ScriptRunContext! This warning can be ignored when running in bare mode.
2025-11-09 17:12:38.404 Thread 'MainThread': missing ScriptRunContext! This warning can be ignored when running in bare mode.
2025-11-09 17:12:38.409 Thread 'MainThread': missing ScriptRunContext! This warning can be ignored when running in bare mode.
2025-11-09 17:12:38.410 Thread 'MainThread': missing ScriptRunContext! This warning can be ignored when running in bare mode.
2025-11-09 17:12:38.410 Thread 'MainThread': missing ScriptRunContext! This warning can be ignored when running in bare mode.
2025-11-09 17:12:38.410 Thread 'MainThread': missing ScriptRunContext! This warning can be ignored when running in bare mode.
c:\Users\mahes\AppData\Local\Programs\Python\Python313\Lib\site-packages\keras\src\saving\saving_lib.py:797: UserWarning: Skipping variable loading for optimizer 'rmsprop', because it has 26 variables whereas the saved optimizer has 50 variables.
    saveable.load_own_variables(weights_store.get(inner_path))
2025-11-09 17:12:38.860 Thread 'MainThread': missing ScriptRunContext! This warning can be ignored when running in bare mode.
2025-11-09 17:12:38.861 Thread 'MainThread': missing ScriptRunContext! This warning can be ignored when running in bare mode.
2025-11-09 17:12:38.861 Thread 'MainThread': missing ScriptRunContext! This warning can be ignored when running in bare mode.
2025-11-09 17:12:38.862 Thread 'MainThread': missing ScriptRunContext! This warning can be ignored when running in bare mode.
2025-11-09 17:12:38.862 Thread 'MainThread': missing ScriptRunContext! This warning can be ignored when running in bare mode.
2025-11-09 17:12:38.863 Thread 'MainThread': missing ScriptRunContext! This warning can be ignored when running in bare mode.
2025-11-09 17:12:38.863 Thread 'MainThread': missing ScriptRunContext! This warning can be ignored when running in bare mode.
2025-11-09 17:12:38.863 Thread 'MainThread': missing ScriptRunContext! This warning can be ignored when running in bare mode.
```